

The New Malthusian Mistake: Why Are Societies Choosing a Shrinking Economy Over True AI Growth?

Maria Elena Mejia. June 1, 2026.

This paper was developed with AI assistance for writing, editing, and research refinement. All arguments, analyses, and conclusions represent the author's original intellectual contribution and responsibility.

It is part of a broader research program on human-centered economics. See: [Human-AI Compact](#)

Abstract

The growing adoption of AI across industries has intensified debates about its consequences for employment, with layoffs increasingly presented as an inevitable outcome of technological progress. This paper argues that such an approach reflects a technical misclassification, positioning the institutional choice of substituting human labor within productivity definitions that rule out managerial responsibility. Drawing on North's institutional economics and the historical parallel of Malthusian static production frontier thinking, this paper reframes the productivity-employment relationship as an outcome shaped by the incentive structures governing how firms, governments, and civil society respond to technological changes. Furthermore, an examination of governance precedents from multiple jurisdictions reveals this capacity in practice. These include the United States Food and Drug Administration (FDA), the European Union Artificial Intelligence Act (EU AI Act), recent Chinese labor court rulings, and California's Executive Order N-6-26 (May 21, 2026) directing assessment of AI workforce impacts. Together, they demonstrate that societies retain both the capacity and the precedent to govern transformative technologies responsibly. This illustrates that protecting workers and sustaining innovation can coexist as complementary goals. The implications of this contribution offer governments, organizations, and civil society a critical analytical perspective for reassessing assumptions about AI deployment and identifying gaps in existing structures regarding workforce protection.

Introduction

Studies claim that millions of workers will face displacement in the coming decade as AI automates white-collar tasks and robots reshape manufacturing floors (Acemoglu & Restrepo, 2019; Acemoglu et al., 2022). Estimates vary significantly: Goldman Sachs projects up to 300 million full-time jobs globally could be affected, while the World Economic Forum forecasts a net displacement of 14 million jobs by 2027. The deeper problem is the narrow purpose to which we are directing AI, namely the substitution of human labor. We keep asking the most basic possible question to the most advanced technology in human history: *How many people can we replace to lower costs?*

From an economic standpoint, reducing staff creates the appearance of efficiency without guaranteeing real productivity growth.. Thus, by framing human workers and AI as competitors for the same economic slices, we are missing the entire point of technology which is to amplify what is possible. More critically, we have elevated AI to a position where it appears to displace workers by itself; when the critical truth is that decision-makers are the architects of those layoffs.

This reductive framing not only misallocates the potential of a transformative technology, but also sets in motion a chain of institutional and macroeconomic consequences that extend far beyond the firm level. Given that almost every economy depends heavily on workers' taxes to fund health, pension, and social programs, if AI displaces a significant share of workers, those systems will collapse. However, emerging legal precedents such as the May 2026 Hangzhou, China court ruling that AI-driven terminations are unlawful and the EU AI Act (2024) are early signals that societies are beginning to recognize this. These developments signal that societies are beginning to build institutional guardrails to prevent possible unjustified AI layoffs. However, such a movement must be massive and widespread; we need frameworks that go further to redefine what success looks like before the damage accumulates in our economies. History, however, offers useful models for how societies have successfully governed transformative technologies before.

The precedent of having institutions that preserve populations' rights already exists; the Food and Drug Administration (FDA) is a clear example. When unsafe food and drugs threatened public health in the early twentieth century, the United States recognized the need to act and established regulatory structures. Technology like social media, AI, and immersive virtual environments deserve a similar level of priority due to the way they reach into daily lives, livelihoods, and mental health of entire populations. These technological innovations have been deployed at scale without the kind of structured assessment the FDA applies to a new drug. After all, the economic and social costs of inaction are distributed across households, communities, and public budgets rather than absorbed by corporations. The consequences for labor markets, mental health, and economic stability are only beginning to surface, and the window for proactive institutional design is narrowing.

The urgency of this moment becomes clearer when examined against a deeper intellectual history. In 1798, Thomas Malthus concluded that population growth would inevitably outpace food supply, leading to a catastrophe of widespread famine, poverty, and suffering. He failed to account for how technology would expand the production frontier entirely. Today, prevailing economic analysis is repeating a similar mistake in its treatment of AI. We are assessing AI's impact on labor using institutional structures of the present, projecting job losses as destiny instead of determining new structures AI could improve. Keynes (1930) saw further; he predicted that technology would deliver 15-hour work weeks within a century, the productivity arrived, but the vision required to distribute its gains equitably never materialized. What Keynes called "technological unemployment" became a foundation for business and economic success, establishing the reduction of labor inputs as one of the primary measures of economic progress. Since then, this blind spot has been in our economies, failing to manage, monitor, and be accountable for the limits of such transformative technology.

It is precisely this accumulated blind spot that the present paper seeks to address, hypothesizing that unemployment outcomes reflect institutional choices about vision and management capability alongside technological change. Institutional arrangements, shaped by incentive structures and governance frameworks, determine the productivity-employment relationship far more than technological capability alone. Since AI's impact stems from institutional choices, we can always intervene, redefine its scope, and align its deployment with the needs of our society. Then, under these circumstances, management must play an active role, making institutional choices, and demanding that all parties take responsibility.

This study therefore addresses the following research question: *What do existing governance precedents reveal about society's capacity to address institutional accountability gaps created by transformative technologies?*

To answer it, this paper proceeds as follows. Section 2 reviews the literature on job displacement, contrarian perspectives, and the foundational theories this paper builds on. Section 3 reframes the productivity-employment relationship as an institutional choice rather than a technological inevitability, examining the analytical distinction between labor substitution and augmentation and the systemic consequences of conflating cost reduction with genuine productivity growth. Section 4 examines existing governance precedents that demonstrate society's capacity to hold institutions accountable for the consequences of transformative technologies, drawing implications for the responsible deployment of AI. Section 5 examines the applied implications and relevance of this analysis across three contrasting economic scenarios, drawing practical insights for governments, organizations, and civil society. The conclusion summarizes the theoretical contribution, acknowledges the study's scope, and proposes directions for future empirical research.

2. Literature Review and Theoretical Foundations

The promise of an automated and futuristic world develops as technology continues evolving, and the proliferation of articles and studies increases alongside it. Although job reduction appears to be the most certain outcome of AI adoption across the economy, there is considerable room for alternative considerations that can enrich the discussion. These considerations center particularly on governance guidelines where AI serves as a tool to enhance human capabilities rather than substitute them. The idea of automation as an inevitable force has prevailed for decades, with some contrarians raising valid concerns, and historical evidence consistently showing that technological advancements have added more work than they have destroyed. During the current AI age, however, we seem to be overestimating what AI fundamentally is, while blindly accepting the majority idea that AI autonomously displaces jobs without interrogating the institutional choices that enable such outcomes.

In economic terms, we are witnessing a bias toward labor substitution, because replacing humans with AI implies reducing costs, which creates cost savings that resemble business improvement but mask the absence of real productivity growth. Labor augmentation, which refers to using technological innovation to expand human output and capacity, has been systematically overlooked, largely because it demands a higher level of critical thinking, visionary leadership, and advanced management skills than simple headcount reduction may require.

Acemoglu and Restrepo (2018, 2019) built robust arguments supporting the idea that AI and robotics could replace workers, effectively establishing loss of employment as the default narrative in economics. One of their central propositions is that decreases in production costs generate a productivity effect that can expand the economy and raise demand for workers in non-automated tasks, creating what they describe as countervailing forces. However, contemporary expectations have pushed toward extreme interpretations of this framework, anticipating massive layoffs as an almost

inevitable outcome. They emphasize that these countervailing effects are incomplete, noting that even when strong, automation increases output per worker more than wages and reduces the share of labor in national income. Their exploration identified that new task creation could counterbalance displacement, yet they identify constraints and imperfections including skill mismatches and excessive automation rates that slow this adjustment. This paper argues that these constraints reflect institutional choices about incentive structures that the following sections will challenge directly, treating what appears as economic law as a fundamentally changeable arrangement.

Autor (2024) challenged the displacement consensus by perceiving AI advances as an opportunity to enhance human expertise across broader worker populations. His framework envisions collaborative workspaces where workers' knowledge and judgment, aligned with AI capabilities, support more sophisticated decision-making than any single actor could achieve independently. Importantly, Autor highlighted a dimension the displacement narrative overlooks, demographic pressures from plummeting birth rates across advanced economies. When properly deployed, AI could address these structural challenges through capability investment alongside technology deployment. Historical precedents such as nuclear fission and electricity generation demonstrate that institutions play a decisive role in determining the purpose and incentives of transformative innovations. He emphasized that the trajectory of AI is ultimately a collective choice involving all actors, one that must weigh income distribution, political influence, and individual rights alongside narrow efficiency measures.

Similarly, Brynjolfsson and McAfee (2017) positioned AI as a general-purpose technology capable of transforming entire industries, while locating the primary bottleneck not in the technology itself but in management, implementation, and organizational imagination. Their argument is that major breakthroughs occur when open-minded and technically capable leaders see past traditional approaches and envision genuinely alternative ones. This suggests that success in AI adoption depends on leadership quality and organizational imagination more than on technological

capability itself. This perspective aligns with the idea that capturing the true gains of AI implementation demands responsive and visionary leadership that is actively seeking opportunities to expand human capability.

A foundational argument reinforcing the displacement narrative rests on the premise that AI is evolving toward full autonomy, capable of independently determining economic outcomes. However, a closer examination of how modern AI systems are architecturally designed reveals significant limitations that undermine this assumption. The core logic of the AI framework relies on *transformer models*, *self-attention mechanisms*, and *pattern recognition* derived entirely from historical data. Transformers belong to a class of deep learning neural network structures that transform long input sequences into output sequences by processing different pieces of a sequence simultaneously, selecting and weighting the most contextually relevant elements to build more robust predictions (Amazon Web Services, 2026). It is precisely this mechanism that has enabled advanced pattern recognition across massive databases, but it is also this same mechanism that structurally prevents AI from exercising genuine agency, forming independent goals, or making autonomous decisions about human labor.

As this architecture makes clear, AI is a system built entirely on past data, which means, it lacks access to future datasets from which to learn labor market futures. Thus, the decision remains an entirely human choice, made by managers, executives, policymakers, and the broader environment in which firms operate. In the absence of demonstrated technological autonomy, which current science provides little evidence for within the near-term horizon, every layoff decision attributed to AI remains in practice a decision made under human authority and therefore subject to institutional accountability. Therefore, displacement discourses that position AI as the primary driver are empirically weak while simultaneously serving to blur responsibility of the actors who ultimately pull the lever.

Standard economic models treat the decision to capture productivity gains through headcount reduction as rational optimization. What these models systematically miss is the formative role that formal and informal norms, meaning the full institutional environment of people, firms, governments, and shared expectations, play in shaping economic behavior over time. North (1990) argued that institutions are the rules of the game in a society, and that economic outcomes are inseparable from the incentive structures those rules create. Similarly, Veblen (1899) identified how entrenched habits of thought within institutions can persist long after the material conditions that produced them have changed. This creates inertia that resists more productive alternatives even when they are clearly available. By assuming that technological development necessarily implies layoffs, frameworks effectively exempt the institutions involved from accountability, allowing a contingent and reversible choice to masquerade as an economic law.

The case of Malthus is particularly instructive for this paper because it evidences how governance and technological structures of a given time can fundamentally limit the vision of even rigorous thinkers. He based his calculations on existing agricultural conditions and arrived at scarcity as a principal result, never accounting for the possibility that technology would expand the production frontier beyond what his models could represent. Today's decision-makers are committing a structurally identical error, projecting forward from existing firm hierarchies, shareholder primacy norms, and quarterly performance incentives, arriving at displacement as the inevitable outcome of AI adoption. Like Malthus, they are treating a historically contingent arrangement as a natural law, and it is precisely that assumption that this paper challenges.

3. Reframing the Productivity-Employment Relationship

This paper argues that the dominant perspective on AI-driven job displacement reflects a biased analytical error, one that confuses a historically contingent arrangement with an economic law. The historical arrangement refers to the norms, incentives, and

organizational habits that societies have deliberately built over time, which can be redesigned when circumstances or priorities change. The economic law, by contrast, implies an inevitable and irreversible outcome, one that removes from view the responsibility that decision-makers hold over the choices they make. These are fundamentally different things, yet the current narrative treats them as equivalent. According to North (1990), the rules governing how productivity gains are distributed are humanly conceived incentive structures, shaped by the governance environment in which firms and governments operate. What the prevailing argument presents as the rational outcome of automation, namely reducing the workforce, is more precisely described as the predictable output of a system that rewards short-term cost minimization over long-term capability investment.

Nevertheless, reframing the relationship between AI and employment as an institutional choice allows us to accept the empirical evidence on automation while situating it within broader incentive structures that shape how firms, governments, and individuals respond to technological change. Understanding how AI can complement rather than replace human capabilities is particularly critical to elevating this conversation from a private business decision to a collective social responsibility. Society itself, operating through both formal regulations and shared expectations about how economies should function, represents the broadest actor in this process, and it is precisely at that level where the standards for human dignity and economic participation must be anchored. Understanding why augmentation remains institutionally unrewarded requires a more careful analytical distinction between the two pathways available to firms and governments when deploying AI.

Labor substitution and augmentation represent two distinct paths that can complement each other when institutions deliberately create the conditions for both to coexist. The former replaces human input with automated processes, reducing the workforce and capturing productivity gains as cost savings. The latter deploys AI to expand human output, enabling workers to perform tasks of greater complexity and value than

previously possible. The prevalence of substitution over augmentation responds to an institutional inertia in Veblen's sense, where habits of thought persist long after the material conditions that produced them have changed, while augmentation requires vision, longer time horizons, and accountability structures that current incentive systems lack mechanisms to reward augmentation approaches. When deployed with intention, substitution of repetitive tasks and augmentation of human skills become complementary forces, with automation handling routine processes while freeing workers to develop capabilities needed for the more complex demands that technological change generates.

Veblen (1899) identified how entrenched habits of thought within institutions persist long after material conditions change. This pattern explains why the familiar path of headcount reduction appears rational when augmentation requires institutional transformation. The case of Germany's Kurzarbeit work-sharing program offers a concrete counter-example. By distributing productivity losses across the workforce through short-time work arrangements rather than concentrating them in layoffs, German institutions demonstrated that augmentation-oriented choices are not only possible but measurably effective, producing faster economic recovery, preserved organizational knowledge, and sustained consumer demand across multiple economic downturns. This case demonstrates that alternatives to substitution-driven approaches are achievable, though implementation varies by context. Nevertheless, this case sets a precedent that refutes the claim that substitution is the only rational institutional response to technological disruption.

A critical distinction requires careful consideration within this analysis. Some level of unemployment is both inevitable and functional for any economy, reflecting the natural adjustments that dynamic labor markets continuously produce. Economists recognize four types: *frictional*, reflecting workers transitioning between positions; *cyclical*, responding to broader economic upturns and downturns; *structural*, resulting from genuine technological obsolescence where retraining may be costly and difficult; and

institutional, caused by the regulations and norms that shape labor market behavior. All four fall outside the scope of this paper's critique. What this paper addresses is a distinct pattern, one where firms deploy AI adoption arguments to justify workforce reductions driven primarily by cost objectives and what they characterize as AI-driven productivity gains. In doing so, they present a discretionary business decision as an inevitable technological outcome, reproducing a *productivity illusion* where cost savings are misrepresented as genuine economic growth.

A central concept in this discussion is productivity, understood in its full technical economic meaning. It has been a subject of paramount importance for economic growth and a sustained concern among foundational economists including Adam Smith, David Ricardo, John Hicks, and Joseph Schumpeter, among others. A desirable outcome for nations and firms alike is to increase it, and historically this has occurred when private and public institutions participate collectively in the process, combining innovation, education, efficiency, and infrastructure investment to expand what economies can produce (World Bank, 2016). Productivity growth, understood at this systemic level, is never the result of a single input or a single decision made in isolation.

When the assessment of productivity is reduced to the firm level, and AI is positioned as the primary justification for workforce reductions, the broader meaning of the concept is significantly narrowed. A central analytical distinction that this interpretation misses is the difference between cost reduction and real productivity growth. From an economics standpoint, true productivity refers to the expansion of output per unit of input, meaning that authentic gains require either more output from the same inputs, or the redeployment of freed resources toward higher-value activities that expand overall economic capacity. When firms reduce their workforce while AI absorbs those tasks and expands output, they are capturing a cost saving while the broader economic system absorbs the social costs of displacement through contracting consumer demand, eroding tax revenues, and increased pressure on public institutions. The difference

between a productivity gain that is broadly distributed and one that is privately concentrated carries significant consequences at the macroeconomic level.

This gap between privately captured cost savings and broadly distributed economic value is what this paper identifies as the productivity illusion, where firm-level efficiency gains mask systemic macroeconomic costs that the broader economy ultimately absorbs.

Brynjolfsson and Hitt (2000) reviewed firm-level evidence demonstrating that investments in information technology are linked to higher productivity when accompanied by organizational transformation. Their central finding is that complementary organizational investments in business processes and work practices are essential for realizing IT productivity gains. Technology and organizational change must occur together to produce measurable results. The organizational choices firms make about how to deploy AI, whether to invest in worker capabilities, restructure work processes, or reduce headcount, are institutional decisions that determine whether productivity gains materialize and how they are distributed. When firms deliberately invest in complementary organizational changes alongside technology adoption, productivity gains expand and distribute more broadly across the workforce and the broader economy. The alternative, where technology deployment prioritizes cost reduction alone, produces gains concentrated among capital owners while labor absorbs adjustment costs. This distinction reveals that AI's labor market impact depends fundamentally on institutional choices about deployment strategy.

The macroeconomic consequences of treating layoffs as a productivity strategy extend far beyond the company level and warrant serious attention. Almost every economy depends heavily on workers as both producers and consumers, meaning that a system systematically reduces employed populations, progressively eroding the consumer demand that sustains those same firms. Tax revenues that fund health, pension, and social programs contract alongside employment, placing additional pressure on public

institutions precisely when displaced workers need them most. Historical episodes of demand collapse, from the Great Depression to the 2008 financial crisis, share a common thread: the concentration of gains among a narrowing group of actors while the broader population loses purchasing power, generating a contractionary spiral that private actors alone cannot reverse. What distinguishes the current moment is that this outcome is not arriving as an external shock but as the foreseeable result of deliberate institutional choices being made daily across firms and industries.

Firms that deploy AI primarily to reduce labor costs may improve their short-term margins, but this approach contributes to a broader macroeconomic dynamic that progressively erodes consumer demand. Without workers earning wages, the revenue base that sustains those same firms ultimately contracts, a dynamic that this paper characterizes as a *productivity illusion with systemic consequences*.

What makes this decision point particularly relevant is that it operates largely without established rules. Unlike financial or environmental decisions that carry regulatory disclosure and legal liability, workforce restructuring driven by AI adoption currently exists in an institutional environment where firms face minimal responsibility for the broader systemic consequences their choices generate. The absence of formal ownership reflects regulatory systems still catching up with the scale and speed of this technological change, rather than a natural feature of market economies. The research question this paper poses, namely what existing governance precedents reveal about society's capacity to address the accountability gap created by AI-driven displacement, is therefore a collective challenge that individuals, organizations, and governments are all entitled and obligated to confront.

Addressing that challenge requires recognizing that the productivity-employment relationship is not a fixed equation but a variable outcome shaped by the environment in which technological decisions are made. Western European economies through active

labor market policies, including job security programs, work-sharing arrangements, and publicly funded retraining schemes, have demonstrated that augmentation-oriented outcomes are achievable without sacrificing competitiveness or growth. Those that have treated workforce reduction as the default response to technological disruption have consistently underestimated the systemic costs that accumulate when consumer demand erodes, tax bases contract, and public institutions absorb the consequences of private decisions made without regard for collective welfare. The distinction between these trajectories is institutional, and it is that distinction that the following section examines through the lens of existing governance precedents.

4. The Case for Institutional Accountability

Since the public release of Large Language Models in 2023, including ChatGPT, Gemini, and others, concerns about the societal consequences of AI deployment have moved across academic, legal, and policy spheres. The scale of adoption across industries has exposed a governance gap that existing frameworks were not designed to address. This raises a concrete concern. AI-driven efficiency gains could translate into mass workforce displacement, impacting economic and social stability. However, our society has faced similar challenges before, and across different historical moments, institutions have recognized that when transformative technologies generate costs borne disproportionately by workers and communities, public accountability becomes necessary. The responses examined below converge on a shared principle: that the costs of technological transformation cannot be unilaterally transferred to workers without oversight and accountability.

The emergence of the US Food and Drug Administration (FDA) in the early twentieth century offers an instructive parallel. Institutional intervention becomes necessary when transformative technologies affect public welfare at scale. Before federal regulation, food and drug manufacturers operated without systematic accountability for the harm

their products caused, and ordinary consumers absorbed those costs directly. The FDA came into existence to ensure that the benefits of innovation reached the public without sacrificing their safety, establishing a precedent that markets operating without oversight consistently shift costs onto the most vulnerable actors. The same logic applies to AI deployment in labor markets. When private actors make decisions that generate broad social costs, like displaced workers, institutional responsibility follows the same rationale that drove the FDA's creation, namely that markets guided solely by their own incentives are unreliable in protecting the public good.

More recent responses reinforce this precedent across different legal and political contexts. In 2024, the European Union approved the Artificial Intelligence Act, establishing the first comprehensive legal framework for AI governance among major economies, reflecting a deliberate judgment that markets alone cannot manage the social consequences of AI deployment. In 2026, the Hangzhou (China) Intermediate People's Court ruled unlawful the dismissal of a quality assurance supervisor whose role had been automated, establishing that a company's decision to adopt AI constitutes a strategic business choice rather than an unforeseeable circumstance. Most recently, California's Executive Order N-6-26 of May 21, 2026 directed state agencies to restructure incentive frameworks around AI deployment and explore worker ownership models, explicitly acknowledging that markets alone will fail in reaching a balanced equilibrium.

Across different legal traditions, economic contexts, and historical moments, these cases reveal that the social costs of technological transformation require accountability structures. Markets guided solely by their own incentives consistently prioritize efficiency over public welfare protection. What accountability demands in the AI era extends beyond legal compliance and into the culture of firms, governments, and civil society. It requires management driven by critical thinking, creative problem solving, and a genuine consideration of public consequences. It demands rigorous and honest analysis of the risks of maintaining the current trajectory, one in which innovators

increasingly dictate the future of work while that future remains a public concern. And it requires a deliberate rewriting of the incentive structures that currently reward substitution over augmentation, so that the institutional environment North described as the rules of the game begins to reflect the full range of actors whose lives those rules govern.

The real societal cost of ignoring job displacement is a preventable economic crisis, one that accountability and responsible innovation can avert if all actors accept their share of responsibility. When employed populations contract systematically, tax revenues that fund health, education, and social programs decline alongside them, placing pressure on public institutions precisely when displaced workers need them most. Consumer demand, which sustains the revenues of the same firms implementing workforce reductions, erodes progressively, generating a contractionary dynamic that no individual firm has the incentive or capacity to reverse alone. Inequality widens as productivity gains concentrate among a narrower group of capital owners, while the broader population absorbs the adjustment costs. These outcomes are predictable institutional failures that follow directly from incentive structures designed without adequate accountability.

5. So What Are the Applied Implications & Relevance?

The analysis developed across the preceding sections establishes two interconnected contributions. First, that the anticipated wave of AI-driven layoffs reflects a technical misclassification, masking a human bias by embedding the choice to substitute labor within technological and productivity definitions that present it as inevitable. Second, that existing governance precedents demonstrate viable alternative paths, suggesting that the harm of irresponsible AI deployment is neither unpredictable nor unmanageable. Translating these contributions into practical terms requires recognizing that no single actor holds sufficient leverage to redirect the current trajectory alone. This

paper also examines three contrasting economic scenarios, ranging from unchecked displacement to augmentation-led transitions and illustrates how the transmission mechanisms differ. A shared commitment across firms, governments, workers, and civil society represents a path to move forward and capture the genuine productive potential of AI without concentrating its costs among the most vulnerable. Whether AI-driven unemployment arrives as a sudden shock or as a gradual erosion, the macroeconomic destination leads to a structurally similar outcome.

The macroeconomic transmission mechanisms of unchecked AI-driven displacement follow a contractionary logic that remains peripheral to productivity discussions, despite its systemic importance. At the firm level, replacing workers with automated systems initially reduces costs and expands output, generating margin improvements that appear rational in isolation. However, as unemployed workers fail to secure alternative positions, household consumption contracts, since displaced workers reduce spending and increase precautionary savings to extend resources through an uncertain and prolonged period of joblessness. Government unemployment support absorbs the initial shock, but if job decreases arrive suddenly and at scale, public budgets face simultaneous revenue contraction and expenditure expansion, creating fiscal pressure precisely when stabilization is most needed. If AI causes unemployment to occur gradually across industries without sufficient labor market absorption, the contractionary pressure distributes across the broader population, generating sustained economic anxiety and reduced consumer confidence even among employed workers who anticipate future job losses.

In either scenario, the aggregate demand contraction eventually reaches the revenues of the same firms that initiated the displacement cycle, producing the self-defeating dynamic this paper characterizes as the productivity illusion.

The alternative transmission mechanism, generated when institutions intervene to ensure that AI deployment reorganizes rather than eliminates workforce capacity,

follows an expansionary logic with measurable macroeconomic benefits. When firms are required to demonstrate genuine business necessity before implementing AI-driven workforce reductions, and when retraining and redeployment obligations accompany automation decisions, displacement becomes temporary rather than permanent. Workers transition through shorter unemployment spells into new roles, maintaining consumption levels that sustain aggregate demand. Tax revenues remain stable, preserving the public capacity that supports the transition. Firms that invest in augmentation rather than substitution develop organizational capabilities and institutional knowledge that pure automation strategies cannot replicate, generating competitive advantages that compound over time. The cases of Germany's Kurzarbeit scheme and Western European active labor market policies demonstrate that this transmission mechanism is not theoretical but historically documented, producing faster recoveries, more resilient demand bases, and more adaptable workforces than economies that defaulted to displacement as the primary adjustment mechanism.

A third scenario, less discussed but increasingly relevant, involves firms that deploy AI augmentation strategically while simultaneously creating entirely new categories of work that did not previously exist. Historical technological transitions, from mechanized agriculture to computerized financial services, consistently generated new occupational categories that absorbed displaced workers over medium-term horizons, provided that frameworks supported the transition period adequately. The current AI transition carries the same potential, but its realization depends on whether incentives reward the organizational imagination required to identify and develop new human roles alongside automated ones. This is precisely where the Malthusian parallel becomes most instructive. The fixed-pie assumption limits the perspective we need to envision. Accountable governance structures, by contrast, enable ideas that distinguish economies that grow through technological transitions from those that merely survive them.

Addressing these dynamics requires each actor to move beyond passive observation and assume a deliberate role in shaping the conditions under which AI deployment occurs. For governments, the practical implication of this analysis calls for a shift toward proactive design. Rather than managing the social costs of displacement after they materialize, as unemployment benefits and retraining programs currently do, the governance precedents examined in Section 4 suggest that intervention is most effective when it establishes accountability structures before displacement occurs at systemic scale. The FDA model is instructive precisely because it created a filter requiring proof of safety before deployment, rather than compensation after harm. Applied to AI labor deployment, this logic supports early warning mechanisms, mandatory impact assessments for large-scale automation decisions, and work-sharing frameworks that distribute adjustment costs across firms and workers rather than concentrating them in layoffs. Society holds private actors responsible for the costs of technological transformation they generate; the institutional architecture built before job losses determines whether those costs become a public crisis or a managed transition.

For organizations, the implication is strategic as much as ethical. Firms that treat AI deployment primarily as a mechanism for workforce reduction are making a locally rational choice that carries globally destructive consequences and, more critically for their own sustainability, a strategically shortsighted one. Workers are simultaneously producers and consumers, and an economic system that systematically reduces the purchasing power of its own labor force eventually contracts the demand base that sustains firm revenues. The cases of companies that have recorded worker processes specifically to replicate and replace them with AI applications represent an extreme version of this logic, one that governance frameworks should categorize as an abuse of informational asymmetry rather than legitimate business optimization. Long-run competitive advantage in an AI-integrated economy will belong to organizations that develop human capabilities alongside technological ones, because the organizational knowledge, judgment, and adaptability that augmentation generates cannot be replicated by substitution strategies that eliminate the human capacity to produce it.

For workers and civil society, the implication is an active rather than passive institutional role. The governance precedents examined in this paper emerged from collective demands that societies placed on the institutions governing them. Workers and civil society organizations are themselves actors in North's framework, capable of reshaping the informal norms and formal rules that determine how AI deployment proceeds. Gradual and transparent introduction of automation, negotiated between firms and workers with adequate transition support, represents a realistic and documented alternative to unilateral displacement decisions. The automation transition is real and its productive potential is genuine, but its social sustainability depends on whether the institutional arrangements surrounding it reflect the interests of the full range of actors it affects, instead of prioritizing only those with the greatest capacity to capture its gains.

Taken together, these implications point toward a single conclusion. An AI-integrated economy definitively influences the future of work, but the collective choices are what ensure a balanced transition.

Conclusion

This paper has argued that the productivity-employment tradeoff is a choice embedded in institutions, consistently misread as a technological law. In doing so, our society repeats the Malthusian error, projecting displacement as an inevitable and unchangeable outcome rather than examining the institutional structures that produce it. The analysis locates AI-driven unemployment at the decision point where productivity gains are distributed, and demonstrates through governance precedents across different legal and historical contexts that alternative outcomes are both achievable and documented. Economics, in this sense, plays an essential role not in designing institutional conditions but in providing the analytical insights that expose

misclassifications, challenge entrenched assumptions, and equip societies with a more honest framework for evaluating the choices they are already making.

For the field of economics, this shift in thinking matters because it restores institutional agency to the center of the analysis, replacing technological determinism with a governing principle that holds identifiable actors responsible for outcomes that have long been treated as beyond human control.

It is worth acknowledging that this study recognizes a natural unemployment rate inherent in any economy. The institutional accountability argument in this paper applies specifically to the space between necessary structural adjustment and opportunistic substitution, a gap that existing frameworks leave largely ungoverned. The choices made under these circumstances shape employment outcomes in ways that are neither inevitable nor beyond accountability. Identifying where that boundary lies represents one of the most critical research directions this paper leaves open.

This perspective finds its clearest historical echo in Keynes. He was right that technology would deliver extraordinary productivity gains. Keynes overlooked how the arrangements surrounding those gains would concentrate their benefits so narrowly while distributing their costs so broadly. This paper is an argument that the vision he carried, of a future where technological progress expands human possibility rather than contracting it, remains available. It arrives through deliberate institutional choice, not as an inevitable consequence of innovation. The economists who matter in this transition are those who help design the institutional conditions under which displacement becomes a choice societies are no longer willing to make.

Given the theoretical perspective this study provides, several directions merit further investigation. An empirical cross-country comparison of active labor market policy

outcomes and displacement rates would allow the institutional accountability argument to be tested against real data, identifying which governance frameworks most effectively redirect productivity gains toward augmentation. A study applying behavioral economics to managerial decision-making would provide a deeper understanding of why labor substitution defaults persist even when augmentation is demonstrably viable, examining the cognitive, incentive-based, and organizational factors that sustain the productivity illusion at the firm level. Finally, a systematic regulatory mapping of emerging AI labor governance structures across economies, building on the cases examined in Section 4, would identify areas of convergence and divergence that could inform more coordinated international institutional responses. Together, these research directions would move the argument of this paper from theoretical contribution toward empirically grounded policy guidance.

References

Acemoglu, D. & Restrepo, P. (2018). Artificial Intelligence, Automation and Work. NBER Working Paper 24196. https://www.nber.org/system/files/working_papers/w24196/w24196.pdf

Acemoglu, D. and Pascual R. (2019). Automation and New Tasks: How Technology Displaces and Reinstates Labor. *Journal of Economic Perspectives* 33 (2): 3–30.
<https://www.aeaweb.org/articles?id=10.1257/jep.33.2.3>

Acemoglu, D. & Restrepo, P. (2020). Robots and Jobs: Evidence from US Labor Markets
<https://shapingwork.mit.edu/wp-content/uploads/2023/10/Robots-and-Jobs-Evidence-from-US-Labor-Markets.p.pdf>

Acemoglu, D., Autor, D., Hazell, J., and Restrepo, P. (2020). AI and Jobs: Evidence from Online Vacancies. NBER Working Paper 28257. <https://doi.org/10.3386/w28257>

Autor, D. (2024). Applying AI to Rebuild Middle Class Jobs. NBER Working Paper 32140 <https://doi.org/10.3386/w32140>

Autor, D., Levy, F. and Murnane, R. (2003). The Skill Content of Recent Technological Change. Quarterly Journal of Economics. 118(4), 1279–1333. <https://ideas.repec.org/a/oup/qjecon/v118y2003i4p1279-1333..html>

Amazon Web Services (2026). What are transformers in artificial intelligence? <https://aws.amazon.com/what-is/transformers-in-artificial-intelligence/#:~:text=Transformers%20are%20a%20type%20of,tracking%20relationships%20between%20sequence%20components.>

Brynjolfsson, E. and Hitt, L. (2000). Beyond Computation: Information Technology, Organizational Transformation and Business Performance. *Journal of Economic Perspectives*, 14(4), 23–48. <https://www.aeaweb.org/articles?id=10.1257%2Fjep.14.4.23>

Brynjolfsson, E. & McAfee, A. (2017). The Business of Artificial Intelligence <https://people.dsv.su.se/~perjons/AI%20and%20Healthcare/The%20Business%20of%20Artificial%20Intelligence.pdf>

EBSCO (2026). U.S. Food and Drug Administration Is Established. <https://www.ebsco.com/research-starters/history/us-food-and-drug-administration-established>

European Union (2024). The Artificial Intelligence Act. <https://aiactinfo.eu/>

Felter, M. (2012). Short-Time Compensation: Is Germany's Success with Kurzarbeit an Answer to U.S. Unemployment? <https://lira.bc.edu/works/publication-article/609td-fme72>

Goldman Sachs Research. (2026). How Will AI Affect the US Labor Market?. Goldman Sachs Economic Research, March 18, 2026. <https://www.goldmansachs.com/insights/articles/how-will-ai-affect-the-us-labor-market>

Hayes, A. (2026). What Is Unemployment? Causes, Types, and Measurement.

<https://www.investopedia.com/terms/u/unemployment.asp>

Kim, Y. E., Loayza, N., & Meza-Cuadra, C. (2026). Productivity as the Key to Economic Growth and Development

<http://documents1.worldbank.org/curated/en/314741472533203058/pdf/108092-BRI-Policy-3.pdf>

Malthus, T. (1798). *An Essay on the Principle of Population*

<https://math.uchicago.edu/~shmuel/Modeling/Malthus.%20An%20essay%20on%20the%20principle%20of%20population.pdf>

North, D. C. (1990). *Institutions, Institutional Change and Economic Performance*. Cambridge University Press.

https://api.pageplace.de/preview/DT0400.9780511251788_A23569047/preview-9780511251788_A23569047.pdf

SCIO (2026). Chinese court defends labor rights in new AI-replacement case.

http://english.scio.gov.cn/chinavoices/2026-04/30/content_118471189.html

The government of California. California Executive Order N-6-26 (2026). *Executive Order on AI Workforce Disruption*.

<https://www.gov.ca.gov/wp-content/uploads/2026/05/5.21.26-AI-Workforce-EO-FINAL-SIGNED.pdf>

Veblen, T. (1899). *The Theory of the Leisure Class: An Economic Study in the Evolution of Institutions*.

Macmillan. https://archive.org/details/theoryofleisurec00vebl_2

Weber, M. (2021). Institutional Unemployment and its Relevance in the Covid-19 Pandemic. Austrian Economics Conference.

https://www.academia.edu/61636369/Institutional_Unemployment_and_its_Relevance_in_the_Covid_19_Pandemic

World Economic Forum. (2025). The Future of Jobs Report 2025.

<https://www.weforum.org/publications/the-future-of-jobs-report-2025/digest/>